

Comparative Analysis of Sequential Models for Mobile Network Traffic Forecasting

Author: Manzi Ivan

Date: September 2026

Dataset: Telecom Italia Milan CDR (SpazioDati / Big Data Challenge)

Models compared: SARIMA (SARIMAX + Fourier terms) · LSTM · TCN

"How do different sequential models compare for one-step-ahead mobile network traffic forecasting, and how does their performance vary across geographical areas with different traffic characteristics?"

Comparative Analysis of Sequential Models for Mobile Network Traffic Forecasting

1. Introduction

Mobile network operators must anticipate short-term fluctuations in cellular traffic to plan capacity, detect anomalies, and allocate radio and backhaul resources proactively rather than reactively. Forecasting this traffic is made difficult by strong daily and weekly periodicity, spatial heterogeneity across geographic areas (dense business districts vs. residential neighborhoods), and occasional bursts driven by events or anomalies. This report investigates **one-step-ahead forecasting of Internet traffic** using the Telecom Italia "Milan Grid" Call Detail Record (CDR) dataset [1]: telecommunications activity across 10,000 geographical grid squares covering Milan, recorded at 10-minute intervals.

Research question: *How do different sequential models compare for one-step-ahead mobile network traffic forecasting, and how does their performance vary across geographical areas with different traffic characteristics?*

The objective is not to declare a single "winning" model but to investigate the problem empirically: to characterize the traffic signal, select and justify three structurally different forecasting approaches informed by prior research and by the data itself, implement them under a common evaluation protocol, and interpret where and why each model succeeds or fails.

A note on data availability. The dataset spans roughly two months (November 2013 – January 2014) and the assignment specifies an evaluation window of December 16–22. The raw archive available for this project contains only the contiguous block November 1–30, 2013 (~10 GB across 30 daily files), plus two disjoint single-day files (December 31, January 1) that were excluded because they cannot support a fair "total traffic over the observation period" ranking or provide contiguous training history. All analysis in this report therefore uses the available November window, with the evaluation week substituted to **November 24–30, 2013** – the last full week of the archive, chosen to mirror the assignment's placement of the evaluation period near the end of the observation window. This substitution is a data-availability constraint, not a methodological choice, and is documented so the pipeline can be re-run unchanged on the December window if the missing files are obtained (only the two date constants in `common.py` need to change).

2. Related Work and Model Selection

Mobile/cellular traffic forecasting has been studied with both classical statistical models and modern machine learning. **SARIMA and its variants** remain competitive baselines for traffic that exhibits strong, stable periodicity, precisely because they explicitly model autocorrelation and seasonal structure with few parameters and are interpretable [2], [3]. However, comparative studies on cellular traffic specifically report that ARIMA-family models are increasingly outperformed by neural

sequence models as more training history becomes available and as the signal exhibits nonlinear, bursty behaviour that a linear model cannot represent [4].

Recurrent architectures, in particular **LSTM** networks, are widely used for cellular/mobile traffic prediction because their gated memory can learn both short-range fluctuations and longer daily/weekly dependencies without hand-engineering lag features [5]. Using two months of Telecom Italia Milan data very similar to the dataset used here, Santos et al. show that recurrent networks trained on raw Internet-activity sequences produce accurate short-term traffic predictions and outperform naive/linear baselines, particularly during periods with irregular demand [6].

Temporal Convolutional Networks (TCN) are a third, structurally distinct family: a stack of dilated causal 1D convolutions with residual connections, processing an entire input window in parallel rather than recurrently [7]. Because dilation grows the receptive field exponentially with depth, a TCN can reach long effective memory with far fewer sequential operations than an RNN, which typically makes training substantially faster; the trade-off is that a TCN's convolutional filters are shared across the whole window and pooled into a fixed-size summary, so — unlike an LSTM's hidden state, which is updated step-by-step and naturally emphasises the most recent input — a TCN has no built-in mechanism for weighting the most recent timestep more heavily than earlier ones, unless the architecture is specifically designed to do so.

Selection rationale, grounded in this project's own exploratory analysis (Section 4): the total-traffic distribution across the 10,000 grid squares is heavily right-skewed (skewness ≈ 4.24 , top 1% of squares carry $\sim 11.3\%$ of total traffic), and the top-traffic squares show strong, stable daily periodicity (confirmed by the ACF/STL analysis in Section 4) punctuated by occasional sharp, short-lived spikes. This combination — *dominant deterministic seasonality* plus *nonlinear, bursty deviations* — is exactly the setting the literature above identifies as informative for comparing a linear/seasonal baseline against nonlinear sequence learners. The three models were therefore selected to be maximally different in how they represent time and nonlinearity, not variations of one architecture:

1. **SARIMA with Fourier (harmonic) seasonal regressors** — an interpretable linear baseline with an explicit seasonal component, serving as the reference against which the value of nonlinearity/ flexibility in the other two models is measured.
2. **LSTM** — a nonlinear, recurrent sequence model expected to capture the bursty, non-repeating residual structure left after removing daily seasonality (visible in the STL residual plot, Section 4), which a linear model cannot.
3. **TCN** — a nonlinear, non-recurrent sequence model that tests whether recurrence is actually necessary to capture that same residual structure, or whether a parallel, convolutional representation of the same history window is sufficient (and, given its architecture, likely faster to train).

Criticism/limitations of each, considered up front: SARIMA assumes a fixed linear relationship and a fixed seasonal shape, so it cannot adapt if a square's daily rhythm changes over the observation period; LSTM requires more data and compute per prediction (sequential, step-by-step backpropagation through time) and its internal representation is opaque; TCN's parallel convolutions are cheap to train but its fixed-size pooled summary of the input window gives it no explicit

mechanism for recency-weighting, which — as Section 6.3 shows — turns out to matter for this signal. These trade-offs are revisited quantitatively in Section 6.

3. Dataset and Data Preparation

3.1 Data and schema

Each daily raw file (`data/raw/sms-call-internet-mi-YYYY-MM-DD.txt`) is a tab-separated, headerless file with one row per `(square_id, time_interval, country_code)`:

```
square_id  timestamp_ms  country_code  sms_in  sms_out  call_in  call_out  internet
```

A given `(square, 10-minute interval)` combination is split across multiple rows — one per active country code — so raw row counts are roughly 20–30× the number of `(square, interval)` pairs actually needed. This project focuses on the `internet` column, matching the assignment's scope. Daily files are 300–400 MB each; the November archive used here totals approximately 10 GB across 30 files.

3.2 Memory management strategy

The target machine has 8 GB of RAM. Loading the full archive naively — `pandas.read_csv` with default settings across all 30 files — is not feasible: a **single day** loaded this way (all 8 columns, default `int64 / float64` dtypes, no aggregation) already peaks at **≈876 MB** resident memory for ~4.84 million rows (measured out-of-process with `/usr/bin/time -l`; see `results/memory_benchmark.csv`). Loading 30 such files while retaining each in memory to concatenate would require tens of gigabytes — well beyond the available RAM.

The implemented strategy (`notebooks/00_data_pipeline.ipynb`) is a **two-pass, chunked, column-pruned, dtype-downcast aggregation**:

- **Column pruning**: only the 2–3 columns actually needed (`square_id`, `internet`, and `timestamp_ms` where required) are read via `usecols`, instead of all 8.
- **Dtype downcasting**: `square_id` as `int32` and `internet` as `float32` instead of pandas' default `int64 / float64`, halving per-value memory.
- **Chunked reading**: each file is streamed in fixed-size chunks (`chunksize=2,000,000` rows) rather than loaded whole, bounding peak memory to the chunk size regardless of file size.
- **Early aggregation**: country codes are collapsed with a `groupby(square_id[, timestamp]).sum()` applied to *each chunk* immediately after reading it, before any concatenation.
- **Two-pass separation of concerns**: Pass 1 computes only the total traffic per square (a 10,000-length accumulator) to rank squares by total traffic, without ever holding a per-timestamp series for all 10,000 squares. Pass 2 re-streams the archive, filtering to the ~5 squares of interest *before* aggregating, so only a few hundred thousand rows are ever materialized.

Applying the same techniques to a single day file reduces peak memory from **≈876 MB to ≈304 MB (a ~65% reduction)**, `results/memory_benchmark.csv`) while collapsing the output from 4.84 million raw rows to the 10,000 needed. Run end-to-end, both passes together process the full 30-day / ~10 GB archive in well under two minutes, with peak memory bounded by chunk size rather than file or archive size throughout [8].

Trade-offs and limitations. The two-pass design trades I/O (the archive is read twice) for memory (neither pass ever holds more than a small accumulator or a filtered subset in RAM); given that combined I/O for this dataset takes well under two minutes, this trade was clearly worthwhile at this scale. This approach would not by itself scale to a use case requiring the full time series for *all* 10,000 squares simultaneously (e.g. city-wide spatio-temporal modelling) — that would require either a true out-of-core engine (Dask/Polars streaming, or a columnar on-disk format such as Parquet) or a machine with substantially more memory. For the single-area forecasting task this assignment requires, though, the two-pass approach is sufficient and keeps the pipeline runnable on a personal laptop.

3.3 Preprocessing per model


Missing 10-minute intervals (rare, from dropped raw rows) are filled via linear interpolation, forward/back-filled at the boundaries (`common.load_square_series`). Model-specific preprocessing (scaling, sequence construction, Fourier terms) is described in Section 5.

4. Exploratory Analysis

4.1 Distribution of total traffic across squares

Over the observation period (Nov 1–30, 2013), total Internet traffic per square is **heavily right-skewed** (skewness ≈ 4.24 , excess kurtosis ≈ 24.3 ; mean $\approx 298,269$ vs. median $\approx 147,153$ — the mean is roughly double the median). The top 1% of squares (100 of 10,000) account for **~11.3%** of total citywide Internet traffic. This is consistent with the well-known spatial heterogeneity of urban mobile demand: a small number of business/transit hubs generate disproportionate traffic relative to the long tail of residential/peripheral squares (see Figure 1, `results/data_summary.csv`).


Figure 1. Distribution of total Internet traffic per square (10,000 squares).



4.2 Top-3 squares and required squares

The three squares with the highest total Internet traffic over the observation period are **5161**, **5059**, and **5259** (totals: 6.33M, 5.83M, and 5.76M respectively — see `data/processed/target_squares.yaml`). Together with the two squares specified by the brief, 4159 and 4556, their first-two-weeks traffic is shown in Figure 2.

Figure 2. Internet traffic, first two weeks of the observation period, for the 3 highest-traffic squares plus the 2 brief-required squares.




All five squares show a clear **daily commuting rhythm** — a trough around 04:00–06:00 and a broad daytime plateau/peak — but differ substantially in *scale*, *peak shape*, and *weekday/weekend contrast*: - **Square 4159** shows the strongest weekday/weekend contrast of the five: visibly lower and flatter peaks on the weekend (Nov 9–10), consistent with a business/office-dominated area where demand is commute-driven. - **Square 4556** has a comparatively "noisier", less sharply peaked profile with a higher baseline overnight, suggesting a more residential or mixed-use area with less pronounced day/night contrast. - **Squares 5059, 5161, and 5259** (the top-3) show the largest peaks and are visibly correlated with each other in timing, consistent with them being geographically adjacent, high-density squares. **Square 5161** shows two large, short-lived spikes (Nov 3 and Nov 10, both Sundays) that are 2–3× its typical Sunday peak — an anomaly worth flagging for the forecasting task.

4.3 Additional analyses on the highest-traffic square (5161)


(a) Autocorrelation (ACF/PACF). Figure 3 shows a strong, slowly-decaying periodic pattern with peaks at lag ≈ 144 and ≈ 288 (i.e. 1 and 2 days at 10-minute resolution) and troughs at half-day offsets — a textbook signature of dominant daily seasonality. The PACF cuts off sharply after 1–2 lags, indicating that, once daily seasonality is accounted for, the remaining short-term dependence is low-order. **Implication for modelling:** this justifies (i) using a *low-order* autoregressive component for SARIMA combined with an explicit seasonal term rather than a large p , and (ii) a sequence length of at least 144 steps (1 day) for the LSTM/TCN so they can see a full seasonal cycle.

Figure 3. ACF and PACF, square 5161.



(b) STL decomposition and stationarity (ADF test). Figure 4 confirms a large, stable daily seasonal component and a slowly-varying trend, with residuals that are small most of the time but show sharp spikes aligned with the anomalous days noted above. The Augmented Dickey-Fuller test on the raw level series rejects the unit-root null (ADF statistic -13.75 , $p \approx 1.1 \times 10^{-25}$; also rejected on the first difference, statistic -9.79 , $p \approx 6.5 \times 10^{-17}$ — see `results/stationarity_tests.csv`). **This is a useful but partial signal:** ADF tests for a stochastic trend, not for seasonality, so a "stationary" verdict here does not mean the series lacks structure to model — it means differencing is *not* strictly required to remove a trend, but the strong deterministic-like daily cycle still needs to be modeled explicitly, which directly motivates the modelling choices in Section 5.

Figure 4. STL decomposition (period=144), square 5161.



5. Methodology

5.1 Forecasting task

Following the assignment's formalization, at each 10-minute interval t in the evaluation week, each model receives the history of `internet` traffic in area a up to and including t , and must produce a one-step-ahead estimate for $t+1$. This is evaluated in a **walk-forward** fashion across the entire evaluation week: predictions are made at every step using the *true* observed history up to that point (not the model's own previous predictions), matching the assignment's formal definition and giving a fair, comparable one-step error for every model.

5.2 Train / evaluation split

For each of the three target squares, the model is trained on **2013-11-01 00:00 → 2013-11-23 23:50** (23 days, 3,312 ten-minute points) and evaluated one-step-ahead across **2013-11-24 00:00 → 2013-11-30 23:50** (7 days, 1,008 points) — the substitute evaluation week described in Section 1.

5.3 Model-specific input representation, preprocessing, and training

SARIMA + Fourier terms. Input: the univariate `internet` series only (no external features). Seasonality is represented with Fourier (sin/cos) regressors at the 144-step (1-day) period — chosen

instead of a full seasonal-ARIMA term (`seasonal_order=(P,D,Q,144)`) because fitting the latter via maximum likelihood at a 144-step seasonal period is computationally impractical to fit and, especially, to tune; Fourier terms give an equivalent smooth seasonal representation at a small fraction of the cost [2, ch. 12] — see `common.py::fourier_terms`. `(p, d, q)` order and the number of harmonics are chosen by grid search (Section 5.4). No additional scaling is applied. At walk-forward time, the model's Kalman-filter state is updated with each newly observed true value via `.append(..., refit=False)` (no re-estimation of coefficients during the evaluation week). The `SARIMAX` wrapper enforces `enforce_stationarity=True, enforce_invertibility=True`, constraining the optimizer to AR/MA roots outside the unit circle; leaving both unconstrained (statsmodels' own example default) is known to let the optimizer converge to explosive AR coefficients on some squares, producing forecasts that diverge to $\pm\infty$ within a few walk-forward steps, so this safeguard is applied unconditionally rather than only when a problem is observed.

LSTM. Input: a sliding window of the last `sequence_length` scaled `internet` values (single feature). Values are scaled with a min-max scaler fit on the training window only (bounds derived only from training data, applied unchanged to the evaluation window, to avoid leakage). Architecture: two stacked LSTM layers with dropout, followed by a single dense output unit; trained with Adam and early stopping on a held-out validation slice. At evaluation time, one-step-ahead predictions use the *true* preceding values (teacher forcing), consistent with the walk-forward protocol.

TCN. Input: the same sliding window as LSTM (single feature, min-max scaled the same way). Architecture: **6 residual blocks** of dilated causal 1D convolutions, dilation rate doubling per block (1, 2, 4, 8, 16, 32), kernel size 3, with dropout after each convolution; each block adds a 1×1 convolution projection on its shortcut connection only where the channel count changes (i.e. only in the first block, where the input's single channel is projected up to the block's filter count). This gives a receptive field of ≈ 253 steps, comfortably covering the 144-step window. The stack is followed by `GlobalAveragePooling1D` (averaging the learned representation across the whole window into a fixed-size vector) and a single dense output unit. Unlike LSTM's or SARIMA's explicit dependence on the *most recent* observation, this pooling step gives every position in the window equal weight by construction — a design property that turns out to matter for the results in Section 6.3. Trained with Adam and the same early-stopping protocol as LSTM.

5.4 Hyperparameter tuning strategy

Given three structurally different models, an iterative, purely manual search across all of them was impractical within the available time budget; **grid search** was used instead, run once on the single highest-traffic square (5161, most representative and highest signal-to-noise), with the last 3 days of the training window held out as a validation slice strictly *before* the evaluation week (so the evaluation week is never touched during tuning). Every combination tried is logged with its parameters, RMSE/MAE/MAPE, and wall-clock time (`results/tuning_results_sarima.csv`, `results/tuning_results_lstm_tcn.csv`). Grids:

- **SARIMA:** $(p,d,q) \in \{(1,1,1), (2,1,2), (3,1,1)\} \times \text{Fourier harmonics} \in \{0, 2, 4\}$.

- **LSTM**: sequence length $\in \{72, 144\} \times$ units $\in \{[32], [64, 32]\} \times$ learning rate $\in \{0.001, 0.0005\}$ (15 tuning epochs with early stopping, patience 3; the winning configuration is retrained for 40 epochs for the final Section 6 results).
- **TCN**: sequence length $\in \{72, 144\} \times$ filters $\in \{16, 32\} \times$ learning rate $\in \{0.001, 0.0005\}$, with the same 15-epoch/early-stopping protocol. Architecture (block count, dilation schedule, kernel size) was fixed by design rather than tuned, so this grid mirrors LSTM's shape (sequence length $\times$ capacity $\times$ learning rate) using `filters` as TCN's capacity knob in place of LSTM's `units`.

6. Results and Discussion

6.1 Hyperparameter tuning outcome


Grid search on square 5161 (Section 5.4) selected: **SARIMA** order `(3,1,1)` with **4** Fourier harmonic pairs; **LSTM** sequence length 144 with two layers of `[64, 32]` units at learning rate 0.001; **TCN** sequence length 144 with 32 filters at learning rate 0.001. Two patterns in the tuning trajectory are worth noting: (i) for SARIMA, adding Fourier harmonics steadily reduced validation RMSE ($\approx 231 \rightarrow \approx 208$ going from 0 to 4 harmonics at order `(1,1,1)`), confirming that the daily seasonal component identified in the EDA carries real predictive value, while the choice of `(p,d,q)` order made comparatively little difference (≈ 208 – 211 RMSE across all three orders at 4 harmonics) — consistent with the EDA's PACF finding of low-order residual dependence once seasonality is accounted for; (ii) for TCN, `sequence_length=144` with `filters=32` clearly beat every `sequence_length=72` configuration and both `filters=16` configurations, suggesting the model benefits from both a full daily window and enough capacity to use it, but even TCN's best validation RMSE (≈ 267) remained higher than every SARIMA and most LSTM configurations tried.

6.2 Per-square performance

Square	Model	MAE	MAPE (%)	RMSE
5161 (top-1)	SARIMA	143.96	12.25	222.86
5161	LSTM	168.73	21.15	237.81
5161	TCN	188.93	18.49	275.80
5059 (top-2)	SARIMA	124.09	10.71	188.13
5059	LSTM	140.54	14.88	199.35
5059	TCN	141.23	12.54	201.71
5259 (top-3)	SARIMA	101.81	9.37	147.61
5259	LSTM	109.84	11.70	153.68
5259	TCN	195.31	25.65	240.50

(Full tables: `results/model_performance_square_{5161,5059,5259}.csv`; cross-square/model bar chart: Figure 5; prediction overlays for all 9 square/model combinations: `figures/pred_sq*_*.png`.)

Figure 5. Model performance across squares and models.




SARIMA wins outright on every square and every metric — a clean sweep, and the most consistent result of the three models. **LSTM is second on every square on MAE/RMSE**, and clearly ahead of TCN on two of the three squares. **TCN is the most variable of the three**: competitive with LSTM on MAPE for the two highest-traffic squares (5161: 18.49 vs. 21.15; 5059: 12.54 vs. 14.88, TCN actually *beating* LSTM's MAPE on both), but by far the worst model on square 5259 across every metric (RMSE 240.50 vs. LSTM's 153.68 and SARIMA's 147.61 — a ~57% gap over LSTM). Timing tells a different story: TCN is consistently the fastest neural model to *train* (~55–75 s per square vs. LSTM's ~105–118 s — see `results/timing_stats.csv`), reflecting its fully parallel (non-recurrent) convolutions, while inference is near-instantaneous for both once trained (well under 1 ms/step). SARIMA trains fastest of all (~3–7 s) but its walk-forward *inference* is by far the most expensive per step (~46–48 ms/step, from the per-step Kalman-filter state update via `.append(refit=False)`) — over 60× slower per step than LSTM/TCN inference, though still fast enough in absolute terms (~47 s for a full 1,008-step evaluation week) not to be a practical concern here.

Why, given the data characteristics from Section 4: the EDA found a signal dominated by a strong, *stable, low-noise* daily cycle (ACF peaking at lag 144, PACF cutting off after 1-2 lags) with comparatively short low-order dependence once that cycle is accounted for. This is close to a best-case scenario for a linear model with an explicit seasonal term: SARIMA does not need to *discover* the seasonal pattern, it is handed the daily Fourier basis directly, leaving only a low-order AR/MA residual structure to estimate. Both LSTM and TCN, in contrast, must *learn* the daily rhythm purely from ~23 days (3,312 steps) of raw sequence, which is a comparatively small dataset for either architecture to reliably learn a 144-step periodic dependency from gradient descent alone. This matches the broader literature finding that neural sequence models' advantage over SARIMA-family baselines typically emerges with *more* training data and *more* irregular/nonlinear signal [4] — neither of which is strongly the case here.

6.3 Failure case: TCN's pooled representation under-tracks high-frequency noise

The clearest, most instructive failure is on **square 5259**, where TCN's RMSE (240.50) is nearly 63% higher than SARIMA's (147.61) — by far the largest model-vs-model gap seen on any square. Comparing the prediction overlays for this square (Figure 6) shows why: SARIMA's forecast tracks even the high-frequency noise in the actual series almost exactly, step to step, while TCN's forecast is visibly **smoother** than the actual signal — it systematically **overshoots during troughs** (predicting $\approx 400\text{--}800$ when actual traffic dips to $\approx 150\text{--}300$ overnight) and slightly **lags/undershoots at sharp peaks**, effectively producing a locally-averaged version of the true trajectory rather than tracking its immediate fluctuations.

Figure 6. Square 5259: TCN (top) vs. SARIMA (bottom) one-step-ahead predictions.



A plausible structural explanation, consistent with the architecture described in Section 5.3: TCN's `GlobalAveragePooling1D` head compresses its entire 144-step receptive field into a single fixed-size vector, weighting every position in the window equally by construction. SARIMA, by contrast, explicitly conditions its next-step forecast on the *most recent* true observation at every walk-forward step (via the Kalman-filter state update), and LSTM's recurrent hidden state is updated sequentially, which naturally emphasises the most recent inputs it has seen. Neither of TCN's comparators has this "equal-weighting" property. Square 5259 is also the lowest-traffic and visibly noisiest of the three top squares in the first-two-weeks plot (Figure 2), so getting the *immediate* local dynamics right matters proportionally more there than on the higher-traffic squares, where the strong seasonal signal dominates the error budget regardless of model. This is consistent with TCN being competitive with

LSTM on MAPE for the two busier, smoother-signal squares (5161, 5059) but falling far behind specifically on the noisier square, where recency-sensitivity — not just seasonal-pattern recognition — is what separates the models.

Practical implication: for an operator choosing between these three models for a heterogeneous set of areas, TCN's training-speed advantage ($\approx 40\%$ faster than LSTM) is real, but its accuracy is not uniformly competitive — specifically, it should not be trusted on lower-traffic, noisier squares without further architectural changes (e.g. replacing or supplementing global average pooling with a mechanism that privileges the most recent timesteps, such as a final recurrent/attention layer or simply using the last-timestep feature vector instead of a pooled average).

7. Conclusion and Future Work

This project investigated one-step-ahead Internet-traffic forecasting on the Milan CDR dataset, comparing a linear/seasonal statistical baseline (SARIMA with Fourier terms), a recurrent deep learning model (LSTM), and a temporal convolutional network (TCN) across the three highest-traffic grid squares. Three findings stand out. First, **the model that best fits the data's actual characteristics beat both neural models, on every square and every metric:** SARIMA, given the daily seasonality identified during EDA directly as a Fourier basis, was the most consistently accurate model and orders of magnitude cheaper to train than either neural network. Second, **TCN's speed advantage over LSTM did not translate into an accuracy advantage:** despite training $\approx 40\%$ faster (parallel convolutions vs. recurrence), TCN's accuracy was mixed — competitive with LSTM on MAPE for the two busier squares but the clear worst performer on the noisiest, lowest-traffic square. Third, **aggregate error metrics can hide a specific, structurally explainable failure mode** — TCN's `GlobalAveragePooling1D` head has no built-in mechanism for weighting recent observations more heavily, which visibly manifested as an over-smoothed forecast that couldn't track square 5259's high-frequency fluctuations, a failure that would not have been obvious from the RMSE table alone.

Limitations. (i) The evaluation window is a substitute for the assignment's specified Dec 16-22 week, due to the raw archive being incomplete (Section 1); the pipeline is fully re-runnable on the correct window by changing two constants in `common.py` once the missing files are obtained. (ii) Hyperparameters were tuned on a single square (5161) and reused across all three for tractability; a per-square tuning pass might close some of the gap between models on squares where their ranking is close. (iii) TCN's architecture (block count, dilation schedule, kernel size) was fixed by design rather than searched; only its capacity (`filters`), input window, and learning rate were tuned, so a broader architecture search — particularly one that addresses the pooling-head limitation identified in Section 6.3 — might change the TCN-vs-LSTM comparison. (iv) SARIMA's numerical fragility (its optimizer can converge to explosive coefficients without the stationarity/invertibility safeguard applied here) means it needs that safeguard to be deployed safely across many areas unattended — a per-square automatic order-search (e.g. `auto_arima` -style) with a validity check would be a more robust production approach than one fixed order. (v) The walk-forward SARIMA evaluation updates the Kalman filter state with true observations every step but never re-estimates its coefficients; over a

longer deployment horizon than one week, periodic re-fitting would likely be needed as the underlying pattern drifts.

Future work. Obtain the true December window to reproduce the assignment's exact evaluation period; extend the comparison to more than the top-3 squares, including low-traffic/residential squares where TCN's pooling-head weakness (Section 6.3) may be even more pronounced; modify the TCN head (e.g. take the last timestep's features instead of pooling, or add an attention layer) to test whether that closes the gap with LSTM on noisy squares; add cross-square (spatial) features, since neighbouring squares in the same transit corridor showed visibly correlated dynamics in the EDA; and evaluate a hybrid approach (e.g. SARIMA residuals modeled by a neural correction term) to combine the seasonal-baseline strength of SARIMA with nonlinear residual-fitting.

References

- [1] G. Barlacchi et al., "A multi-source dataset of urban life in the city of Milan and the Province of Trentino," *Scientific Data*, vol. 2, p. 150055, 2015.
- [2] R. J. Hyndman and G. Athanasopoulos, *Forecasting: Principles and Practice*, 2nd ed. OTexts, 2018.
- [3] G. E. P. Box, G. M. Jenkins, G. C. Reinsel, and G. M. Ljung, *Time Series Analysis: Forecasting and Control*, 5th ed. John Wiley & Sons, 2015.
- [4] A. Azari, P. Papapetrou, S. Denic, and G. Peters, "Cellular Traffic Prediction and Classification: A Comparative Evaluation of LSTM and ARIMA," in *Discovery Science (DS 2019)*, LNCS vol. 11828, Springer, 2019, pp. 129–144.
- [5] S. Hochreiter and J. Schmidhuber, "Long Short-Term Memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [6] G. L. Santos, P. Rosati, T. Lynn, J. Kelner, D. Sadok, and P. T. Endo, "Predicting short-term mobile Internet traffic from Internet activity using recurrent neural networks," *International Journal of Network Management*, 2021.
- [7] S. Bai, J. Z. Kolter, and V. Koltun, "An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling," *arXiv preprint arXiv:1803.01271*, 2018.
- [8] W. McKinney, "Data Structures for Statistical Computing in Python," in *Proc. 9th Python in Science Conference*, 2010, pp. 56–61.

Full BibTeX entries: `report/references.bib`.